

Nordic probabilistic AI school

Introduction to probabilistic programming languages (PPLs)

Helge Langseth, Andrés Masegosa, and Thomas Dyhre Nielsen

June 14, 2022

Day 1: Probabilistic programming


- Introduction to probabilistic programming
- Probabilistic programming in Pyro

Day 2: Variational inference

- Recap of variational inference (variational inference as optimization)
- Derivation and implementation of selected examples
 - Bayesian linear regression
 - Factor analysis
 - ...

Day 3: Variational inference – cont'd


- Black box variational inference
- Variational inference in Pyro
- Variational auto-encoders



The development of **machine learning systems** requires enormous efforts.



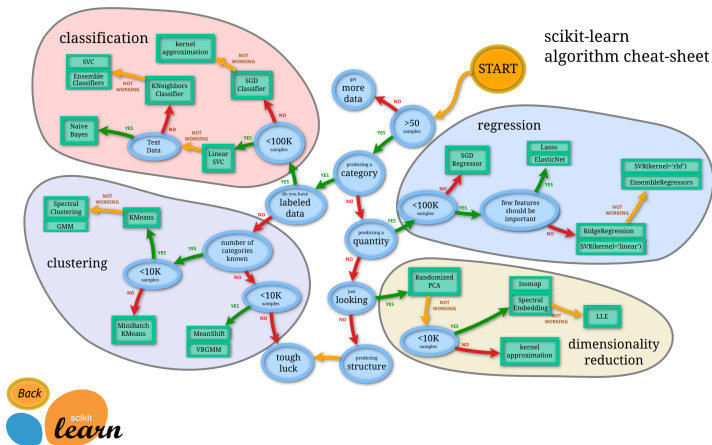
Data Science



Copyright © 2014 by Steven Geringer Raleigh, NC.
Permission is granted to use, distribute, or modify this image,
provided that this copyright notice remains intact.

The development of **machine learning systems** requires enormous efforts.

- It requires of highly qualified experts.




The development of machine learning systems requires enormous effort.

- It is necessary to have highly qualified experts.
- It is difficult to find the ML model most suitable for an application.

Hidden Technical Debt in Machine Learning Systems

D. Sculley, Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips
{dsculley, gholt, dgg, edavydov, toddphillips}@google.com
Google, Inc.




The development of machine learning applications requires enormous effort.

- It is necessary to have highly qualified experts.
- It is difficult to find the ML model most suitable for an application.
- **Programming a ML model is a complex task where many problems are intermingled.**

Wanted: Artificial intelligence experts

In artificial intelligence, job openings are rising faster than job seekers.




Consequences:

- Shortage of AI experts (and high salaries).


Wanted: Artificial intelligence experts

In artificial intelligence, job openings are rising faster than job seekers.

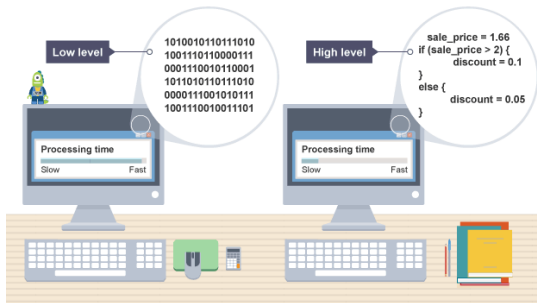


Consequences:

- Shortage of AI experts (and high salaries).
- Only big corporations have the resources for developing ML systems.




Similar situation than 50 years ago:




Similar situation than 50 years ago:

- People used to program in low-level programming languages.



Similar situation than 50 years ago:


- People used to program in low-level programming languages.
- Programming was complex and demand high-expertise.




Similar situation than 50 years ago:

- People used to program in low-level programming languages.
- Programming was complex and demand high-expertise.
- Focus on application and low-level hardware details.

Programming Languages




High-level programming languages brought many advantages:




High-level programming languages brought many advantages:

- Programmers focused on the applications.




High-level programming languages brought many advantages:

- Programmers focused on the applications.
- Hardware Experts focused on compilers.



High-level programming languages brought many advantages:

- Programmers focused on the applications.
- Hardware Experts focused on compilers.
- High gains in productivity.



High-level programming languages brought many advantages:

- Programmers focused on the applications.
- Hardware Experts focused on compilers.
- High gains in productivity.
- “Democratization” of the software development.



Claire D. Costa. Best Python Libraries for Machine Learning and Deep Learning.

<https://towardsdatascience.com/best-python-libraries-for-machine-learning-and-deep-learning-b0bd40c7e8c>

Big Data and Machine Learning Libraries:

- **High-quality**, well-maintained and open-source libraries



Claire D. Costa. Best Python Libraries for Machine Learning and Deep Learning.

<https://towardsdatascience.com/best-python-libraries-for-machine-learning-and-deep-learning-b0bd40c7e8c>

Big Data and Machine Learning Libraries:

- **High-quality**, well-maintained and open-source libraries
- They try to provide **high-level abstractions**.



Claire D. Costa. Best Python Libraries for Machine Learning and Deep Learning.

<https://towardsdatascience.com/best-python-libraries-for-machine-learning-and-deep-learning-b0bd40c7e8c>

Big Data and Machine Learning Libraries:

- **High-quality**, well-maintained and open-source libraries
- They try to provide **high-level abstractions**.
- Hiding under the hood **low level details**.



Claire D. Costa. Best Python Libraries for Machine Learning and Deep Learning.


<https://towardsdatascience.com/best-python-libraries-for-machine-learning-and-deep-learning-b0bd40c7e8c>

Big Data and Machine Learning Libraries:

- **High-quality**, well-maintained and open-source libraries
- They try to provide **high-level abstractions**.
- Hiding under the hood **low level details**.
- Increase the **adoption** of these technologies.

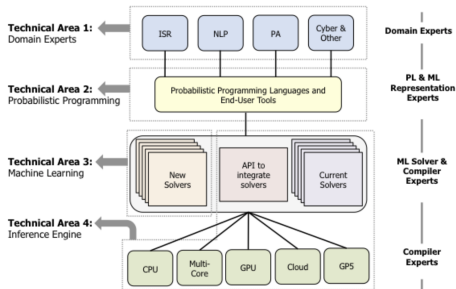
Which are the "high-level libraries" in Probabilistic AI?

Why PPLs?



PPLs as high-level programming languages for **probabilistic machine learning systems**:


- Stacked architecture



PPLs as high-level programming languages for **probabilistic machine learning systems**:

- Stacked architecture
- Different Domain Experts will code their models using the same language.


Why PPLs?



PPLs as high-level programming languages for **probabilistic machine learning systems**:

- Stacked architecture
- Different Domain Experts will code their models using the same language.
- ML experts will focus on the development of new ML solvers.


Why PPLs?



PPLs as high-level programming languages for **probabilistic machine learning systems**:

- Stacked architecture
- Different Domain Experts will code their models using the same language.
- ML experts will focus on the development of new ML solvers.
- Compile experts will focus on running these ML solvers on specialized hardware.

Box's Loop




[Box, 1980; Rubin, 1984; Gelman+ 1996; Blei, 2014]

Benefits of PPLs for developing probabilistic machine learning systems:

- Simplify probabilistic machine learning model code.

Box's Loop




[Box, 1980; Rubin, 1984; Gelman+ 1996; Blei, 2014]

Benefits of PPLs for developing probabilistic machine learning systems:

- Simplify probabilistic machine learning model code.
- Reduce development time and cost to encourage experimentation.

Box's Loop




[Box, 1980; Rubin, 1984; Gelman+ 1996; Blei, 2014]

Benefits of PPLs for developing probabilistic machine learning systems:

- Simplify probabilistic machine learning model code.
- Reduce development time and cost to encourage experimentation.
- Reduce the necessary level of expertise.


Box's Loop



[Box, 1980; Rubin, 1984; Gelman+ 1996; Blei, 2014]


Benefits of PPLs for developing probabilistic machine learning systems:

- Simplify probabilistic machine learning model code.
- Reduce development time and cost to encourage experimentation.
- Reduce the necessary level of expertise.
- “Democratization” of the development of probabilistic ML systems.




1st Generation of PPLs :

- Bugs, WinBugs, Jags, Figaro, etc.




1st Generation of PPLs :

- Bugs, WinBugs, Jags, Figaro, etc.
- Turing-complete probabilistic programming languages. (i.e. they can represent any computable probability distribution).




1st Generation of PPLs :

- Bugs, WinBugs, Jags, Figaro, etc.
- Turing-complete probabilistic programming languages. (i.e. they can represent any computable probability distribution).
- Inference engine based on Monte Carlo methods.




1st Generation of PPLs :

- Bugs, WinBugs, Jags, Figaro, etc.
- Turing-complete probabilistic programming languages. (i.e. they can represent any computable probability distribution).
- Inference engine based on Monte Carlo methods.
- They did not scale to large data samples/high-dimensional models.




2nd Generation of PPLs :

- Infer.net, Factorie, Amidst, etc.




2nd Generation of PPLs :

- Infer.net, Factorie, Amidst, etc.
- Inference engine based on message passage algorithms and/or variational inference methods.



2nd Generation of PPLs :

- Infer.net, Factorie, Amidst, etc.
- Inference engine based on message passage algorithms and/or variational inference methods.
- They did scale to large data samples/high-dimensional models.



2nd Generation of PPLs :

- Infer.net, Factorie, Amidst, etc.
- Inference engine based on message passage algorithms and/or variational inference methods.
- They did scale to large data samples/high-dimensional models.
- Restricted probabilistic model family (i.e. factor graphs, conjugate exponential family, etc.)



PYMC3



3rd Generation of PPLs :

- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.



PYMC3



3rd Generation of PPLs :

- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.
- Black Box Variational Inference and Hamiltonian Monte-Carlo.



PYMC3



3rd Generation of PPLs :

- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.
- Black Box Variational Inference and Hamiltonian Monte-Carlo.
- They did scale to large data samples/high-dimensional models.



PyMC3



3rd Generation of PPLs :

- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.
- Black Box Variational Inference and Hamiltonian Monte-Carlo.
- They did scale to large data samples/high-dimensional models.
- Turing-complete probabilistic programming languages.



PyMC3



3rd Generation of PPLs :

- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.
- Black Box Variational Inference and Hamiltonian Monte-Carlo.
- They did scale to large data samples/high-dimensional models.
- Turing-complete probabilistic programming languages.
- Rely on deep learning frameworks (TensorFlow, Pytorch, Theano, etc).



PyMC3



3rd Generation of PPLs :

- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.
- Black Box Variational Inference and Hamiltonian Monte-Carlo.
- They did scale to large data samples/high-dimensional models.
- Turing-complete probabilistic programming languages.
- Rely on deep learning frameworks (TensorFlow, Pytorch, Theano, etc).
 - Specialized hardware like GPUs, TPUs, etc.





PyMC3



3rd Generation of PPLs :


- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.
- Black Box Variational Inference and Hamiltonian Monte-Carlo.
- They did scale to large data samples/high-dimensional models.
- Turing-complete probabilistic programming languages.
- Rely on deep learning frameworks (TensorFlow, Pytorch, Theano, etc).
 - Specialized hardware like GPUs, TPUs, etc.
 - Automatic differentiation methods.






Pyro's main features (www.pyro.ai) :

- Initially developed by UBER (the car riding company).
- Actively maintained by community contributors (there is dedicated team at Broad Institute)




Pyro's main features (www.pyro.ai) :

- Initially developed by UBER (the car riding company).
- Actively maintained by community contributors (there is dedicated team at Broad Institute)
- Rely on Pytorch (Deep Learning Framework).



Pyro's main features (www.pyro.ai) :

- Initially developed by UBER (the car riding company).
- Actively maintained by community contributors (there is dedicated team at Broad Institute)
- Rely on Pytorch (Deep Learning Framework).
- Enable GPU acceleration and distributed learning.



Pyro's main features (www.pyro.ai) :

- Initially developed by UBER (the car riding company).
- Actively maintained by community contributors (there is dedicated team at Broad Institute)
- Rely on Pytorch (Deep Learning Framework).
- Enable GPU acceleration and distributed learning.

<https://github.com/PGM-Lab/2022-ProbAI>